

# Cache Memory – Part I

**Dr. B. R. Bhowmik**  
**Dept. of CSE**  
**NIT Karnataka**

# Cache Basics

- Cache memory is intended to give memory speed approaching that of the fastest memories available, and at the same time provide a large memory size at the price of less expensive types of semiconductor memories. Thus, cache
  - Is a small amount of fast memory
  - Sits between normal main memory and CPU
  - May be located on CPU chip or module
- The cache contains a copy of portions of main memory.
- When the processor attempts to read a word of memory, a check is made to determine if the word is in the cache.
- If so, the word is delivered to the processor.
- If not, a block of main memory, consisting of some fixed number of words, is read into the cache and then the word is delivered to the processor.
- Because of the phenomenon of locality of reference, when a block of data is fetched into the cache to satisfy a single memory reference, it is likely that there will be future references to that same memory location or to other words in the block.

# Overview of a Basic Cache




Figure 1: Cache overview

- Saves a **subset** of **memory values**
  - We can either have **hit** or **miss**
  - The **load/store** is **successful** if we have a **hit**

# Cache Hierarchy





Figure 2: Cache hierarchy

- Cache uses a **hierarchical memory** system
- L1 (SRAM cells), L2 (SRAM cells), Main Memory (DRAM cells)


# Cache Hierarchy Contd...

- Cache Hierarchy
  - The main memory contains all the memory locations
  - The caches contain a subset of memory locations
- Single and multi-level cache organization with main memory are shown in Figure 3.
- Fig. 3(a) depicts the use of single level of cache.
- Fig. 3(b) depicts the use of multiple levels of cache.
- The L2 cache is slower and typically larger than the L1 cache, and the L3 cache is slower and typically larger than the L2 cache.
- The L1 cache has size 8-64 KB and is composed of SRAM cells.
- The L2 cache has size 128 KB – 4 MB and is also made of SRAM cells.
- The main memory is much larger with size 1 – 64 GB and made of DRAM cells.

# Cache Organization Diagram



(a) Single cache



(b) Three-level cache organization

Figure 3: Cache and main memory relation.

# Cache and Main Memory

- Figure 4 depicts the structure of a cache/main-memory system.
- Main memory consists of up to  $2^n$  addressable words, with each word having a unique  $n$ -bit address.
- For mapping purposes, this memory is considered to consist of a number of fixed length blocks of  $K$  words each. That is, there are  $M = 2^n/K$  blocks in main memory.
- The cache consists of  $m$  blocks, called **lines**.
- Each line contains  $K$  words, plus a tag of a few bits. Each line also includes control bits (not shown), such as a bit to indicate whether the line has been modified since being loaded into the cache.
- The length of a line, not including tag and control bits, is the **line size**. The line size may be as small as 32 bits, with each “word” being a single byte; in this case the line size is 4 bytes.

# Cache/Main Memory Structure





Figure 4. cache and main memory structure





# Cache operation – overview

- CPU requests contents of memory location
- Check cache for this data
- If present, get from cache (fast)
- If not present, read required block from main memory to cache
- Then deliver from cache to CPU
- Cache includes tags to identify which block of main memory is in each cache slot

# Basic Cache Operations

- **lookup** → Check if the memory location is present
- **data read** → read data from the cache
- **data write** → write data to the cache
- **insert** → insert a block into a cache
- **replace** → find a candidate for replacement
- **evict** → throw a block out of the cache

# Cache Lookup

A. Let us a cache have **two SRAM arrays (Figure 5)**

- **tag array** → Saves a part of the block address such that the block can be uniquely identified
- **block array** → Saves the contents of the block
- Both the arrays have the same number of entries
- Efficient cache lookup means to find out a block efficiently in
  - Fully Associative Cache
  - Direct Mapped Cache
  - Set Associative Cache

B. **Running example** : 8 KB Cache, block size of 64 bytes, 32 bit memory system.

# Structure of a Cache


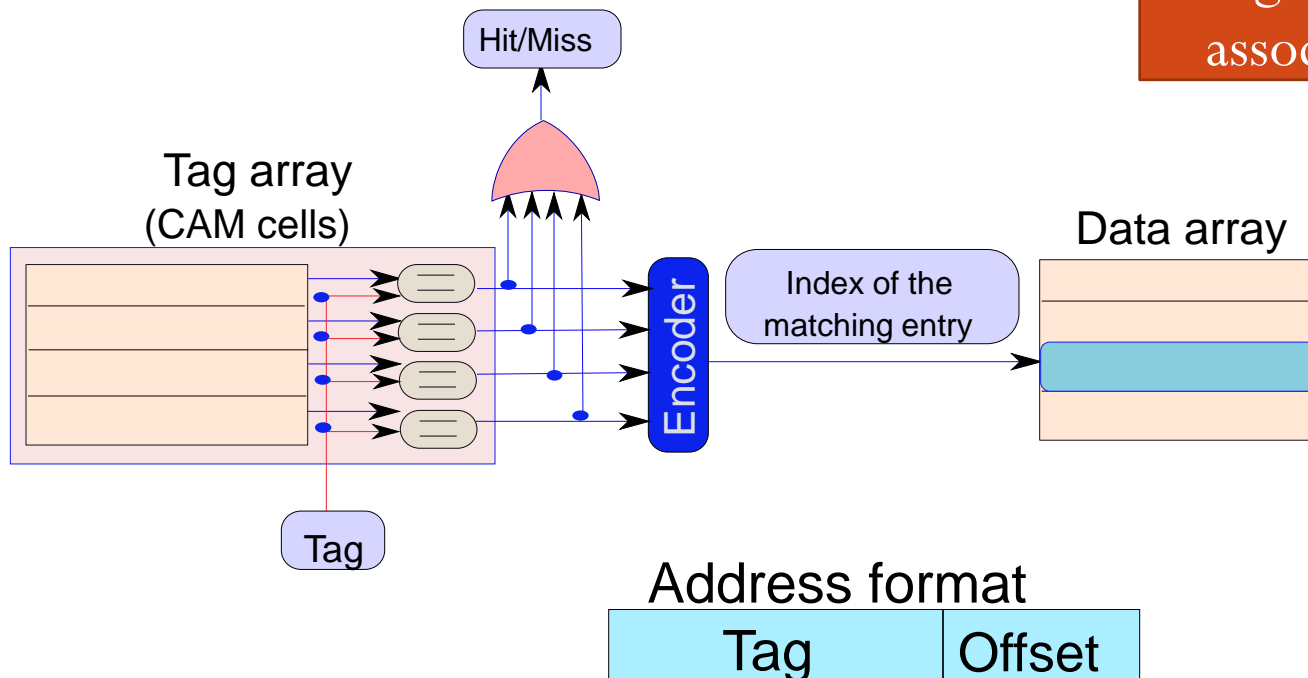


Figure 5. Structure of Cache with Tag and data Array.

# Fully Associative Cache

- We have  $2^{13} / 2^6 = 128$  entries
- A **block** can be saved in any entry
- **26 bit tag**, and **6 bit offset**

Figure 6. A fully associative cache



# Implementation of the FA Cache

- We use an array of CAM cells for the tag array
- Each entry compares its contents with the tag
- Sets the match line to 1
- The OR gate computes a hit or miss
- The encoder computes the index of the matching entry.
- We then read the contents of the matching entry from the block array

# Direct Mapped Cache

- Each **block** can be **mapped** to only **1 entry**

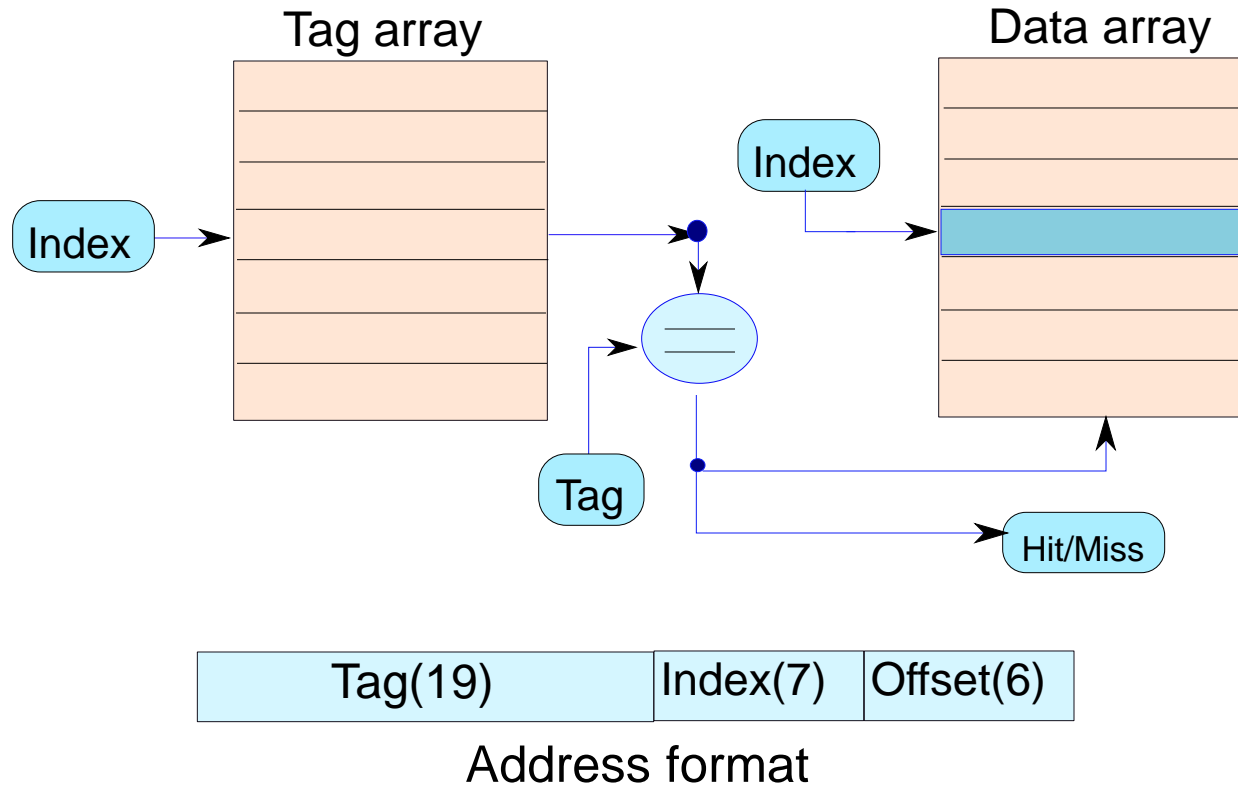



Figure 7. Direct mapped cache

# Direct Mapped Cache Contd...

- We have 128 entries in our **cache**.
- We compute the **index** as  $idx = \text{block address} \% 128$
- We access **entry**, **idx**, in the **tag array** and compare the contents of the **tag** (19 msb bits of the **address**)
- If there is a **match**  $\rightarrow$  **hit** else  $\rightarrow$  **miss**
- Need a solution that is in the middle of the spectrum



# Set Associative Cache



- Let us assume that an **address** can **reside** in 4 **locations**
  - Access all 4 locations, and see if there is a hit
- Thus, we have  $128/4 = 32$  **indices**
- Each **index** points to a *set* of 4 **entries**
- → We now use a 21 bit **tag**, 5 bit **index**

# Set Associative Cache Contd...

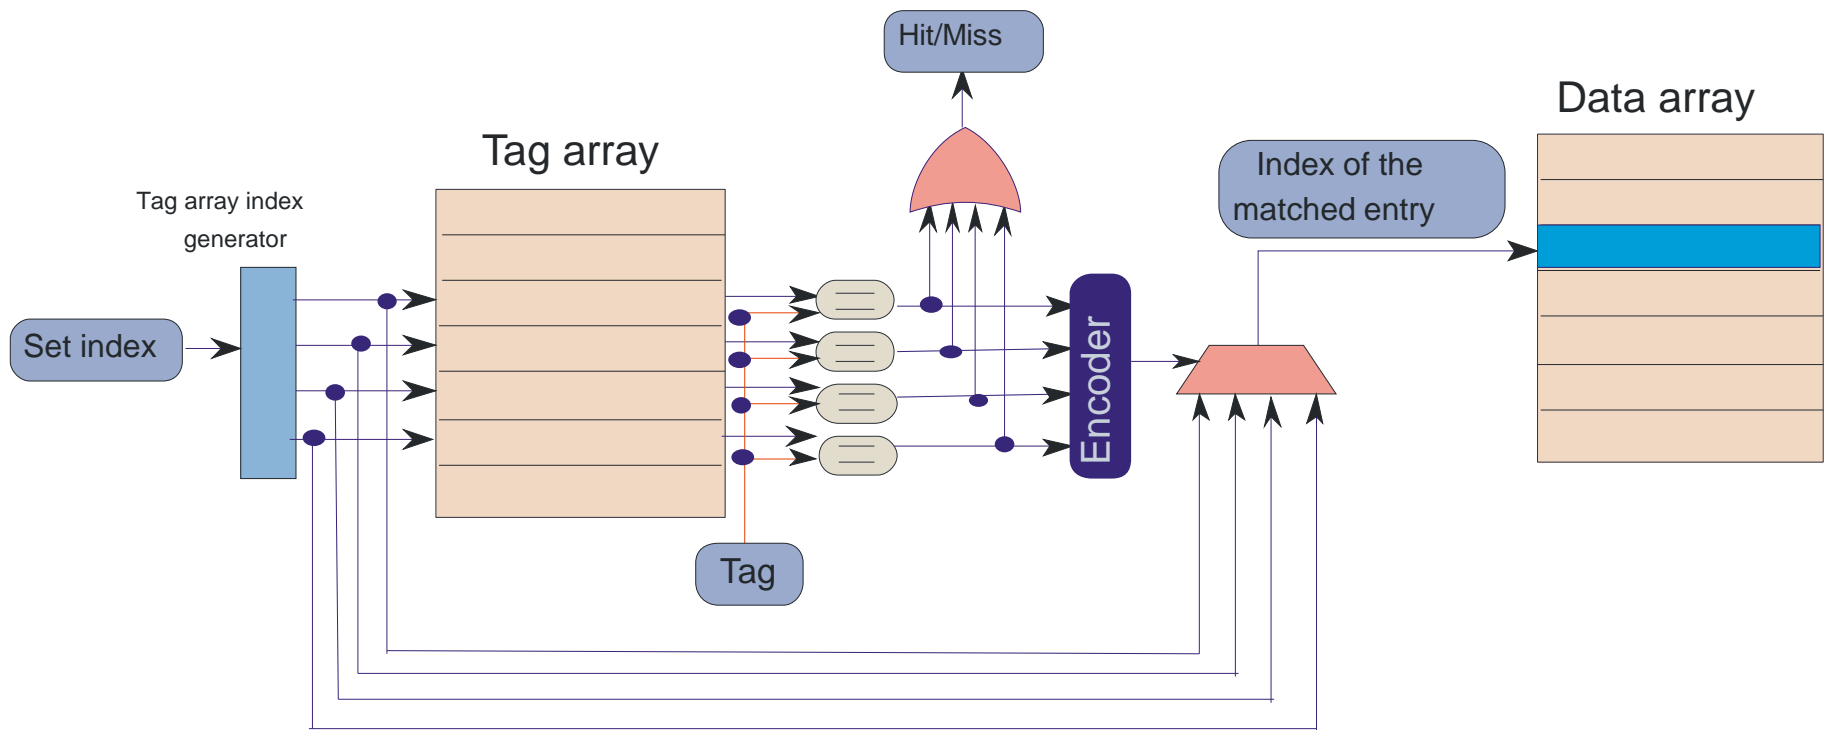


Figure 8. Set associative cache

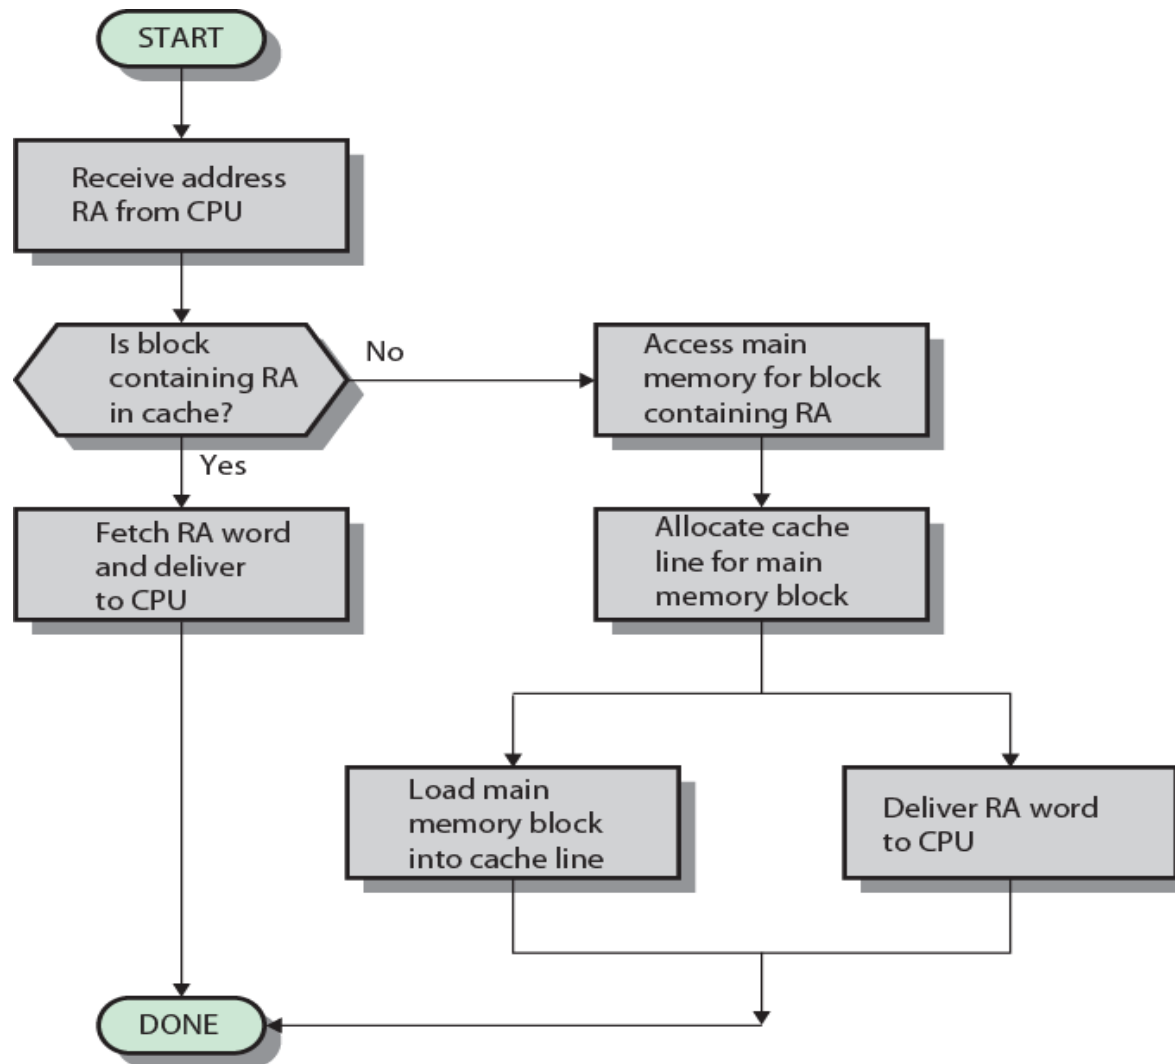
# Set Associative Cache Contd...

- \* Let the **index** be  $i$ , and the **number of elements** in a set be  $k$ 
  - \* We access **indices**,  $i*k$ ,  $i*k+1$ , ...,  $i*k + (k-1)$
  - \* Read all the **tags** in the set
  - \* **Compare** the tags with the tag obtained from the address
  - \* Use an **OR** gate to compute a **hit/ miss**
  - \* Use an **encoder** to find the **index** of the **matched entry**

# Data read operation

- This is a **regular** SRAM access.
- Note that the **data read** and **lookup** can be overlapped for a **load access**
  - We can issue a parallel **data read** to all the ways in the cache
  - Once, we compute the index of the **matching** tag, we can choose the **correct** result with a **multiplexer**.

# Cache Read Operation - Flowchart




# Data write operation

- Before we **write** a **value**, we need to **ensure** that the block is present in the **cache**.
- Why ?
  - \* Otherwise, we have to **maintain** the indices of the bytes that were written to
  - \* We treat a **block** as an atomic unit
  - \* Hence, on a **miss**, we fetch the entire **block** first
- Once a block is there in the cache, Go ahead and write to it .
- Maintain a **modified** bit in the tag array.

# Data write operation Contd...

- If a **block** has been written to, after it was **fetch**ed, set it to 1.



- Data write operation follows write policies.
  - Write through
  - Write back

# Write Policies


- **Write through** → Whenever we write to a cache, we also write to its **lower level**
  - Advantage : Can **seamlessly evict** data from the cache
- **Write back** → We do not **write** to the lower level. Whenever we write, we set the **modified bit**.
  - At the time of **eviction** of the line, we check the value of the **modified bit**



# Insert Operation

If we don't find a block in a cache. We fetch it from the lower level. Then we insert the block in the cache → insert operation

- Let us add a **valid** bit to a **tag**
  - If the line is non-empty, valid bit is 1
  - Else it is 0
- Structure of a tag



# *Insert Operation Contd...*

- Check if any way in a set has an **invalid** line
  - If there is one, then write the **fetch** line to that **location**, set the **valid** bit to 1.
- Otherwise, find a candidate for **replacement**.

# The replace operation

- A cache **replacement** scheme or **replacement** policy is a method to **replace** an entry in the set by a new entry.
- **Replacement Schemes**
  - FIFO Replacement Scheme
  - Least Recently Used
  - Psuedo-LRU
  - Random Replacement Scheme

# Replacement Schemes: FIFO

- For replacement, choose the way with the highest counter (oldest).
- Problems :
  - Can **violate** the principle of temporal locality
  - A line fetched early might be accessed very frequently.

# LRU (least recently used)

- A. **Replace** the **block** that has been accessed the least in the **recent past and** Most likely we will not access it in the near future.
- Directly follows from the definition of **stack distance**
  - Sadly, we need to do more **work** per access
  - Proved to be **optimal** in some restrictive scenarios
- B. **True** LRU requires saving a hefty **timestamp** with every way
- Let us implement pseudo-LRU

# Pseudo-LRU

- A. Let us try to **mark** the **most recently used** (MRU) elements.
- B. Let us **associate** a 3 bit counter with every **way**.
- Whenever we **access** a **line**, we **increment** the counter.
  - We stop **incrementing** beyond 7.
  - We periodically **decrement** all the **counters** in a set by 1.
  - Set the counter to 7 for a newly **fetch**ed block
  - For **replacement**, choose the **block** with the **smallest** counter.

# evict Operation

- If the cache is **write-through**
  - Nothing needs to be done
- If the cache is **write-back**
  - AND the **modified** bit is 1
  - Write the line to the **lower level**

**Thank You**